

# Scaling Language Models

Antoine Bosselut

# Announcements

- **No Lecture Tomorrow!**
- **Assignment 3:** Due Sunday, 21/04/2024 at 11:59 PM
  - Office Hours: **tomorrow**, Thursday 2 PM
- **Assignment 1:** Grading Feedback Session Tomorrow at 2 PM
- **Course Project:** Kickoff!
  - Data Packages were released. If you haven't received them yet, contact us.

# Next few days: Project Sign-ups

- **To-Dos:**


- **URGENT:** Look over the Project Description
- **URGENT:** Sign up for project repository
- Look through README in project repository for details on milestone submission
- Get started early! **Milestone 1 due May 5th!**

# Today's Outline

- **Lecture**

- **Quick Recap:** Scale
- **Managing scale when training:** Scaling laws
- **Managing scale when deploying:** Model Compression (pre-recorded video)
  - how can we make LLMs more compute- and memory-efficient for deployment ?

# Language Model Scaling



Larger models

More data

More compute

More




# Every part of the model scales!

Model Name	$n_{\text{params}}$	$n_{\text{layers}}$	$d_{\text{model}}$	$n_{\text{heads}}$	$d_{\text{head}}$	Batch Size	Learning Rate
GPT-3 Small	125M	12	768	12	64	0.5M	$6.0 \times 10^{-4}$
GPT-3 Medium	350M	24	1024	16	64	0.5M	$3.0 \times 10^{-4}$
GPT-3 Large	760M	24	1536	16	96	0.5M	$2.5 \times 10^{-4}$
GPT-3 XL	1.3B	24	2048	24	128	1M	$2.0 \times 10^{-4}$
GPT-3 2.7B	2.7B	32	2560	32	80	1M	$1.6 \times 10^{-4}$
GPT-3 6.7B	6.7B	32	4096	32	128	2M	$1.2 \times 10^{-4}$
GPT-3 13B	13.0B	40	5140	40	128	2M	$1.0 \times 10^{-4}$
GPT-3 175B or “GPT-3”	175.0B	96	12288	96	128	3.2M	$0.6 \times 10^{-4}$

- Trained on 570 GB of Common Crawl data
- **How?** Used cluster provided by Microsoft



# Why scale?



- Last week, we talked about benefits of scaling in terms of **emergence**
- Practically, training for longer also leads to lower test loss
- Larger models can reach lower test losses


# What should we scale?

Model size, dataset size, compute budget

Given a compute budget, how big of a model can we train?  
and how big of a dataset should we train it on?



# Impact of compute budget




**Dotted lines estimate these curves.  
Need to predict for larger models!**

- For a fixed compute budget, there is an optimal number of parameters that we can train
- Having **too large** a model for **too small** a compute budget does NOT let the model learn
  - Model doesn't see enough examples during training
- Having **too small** a model for too large a compute budget is also bad
  - Repeated computation isn't helpful if the model has no capacity to encode additional information

# Consideration


- With a fixed compute budget (in FLOP-days), we have two costs:
  - Number of floating point operations needed to train on a single example (model size)
  - Number of total examples we will train on
- **How should we trade off these two costs?**
  - Which should we prioritise?

# Model-Data Trade-off




- Larger models benefit more from larger datasets
- **Smaller models saturate**
  - Only so much capacity to learn!
- At some point, larger models don't benefit more from same-sized data
- Model size needs to be scaled jointly with data size

# Other Cool Findings




No need to make models terribly deep



Multiple ratios of depth vs. width (aka embedding size) are fine




# Other cool findings



- LSTMs also follow scaling laws, benefitting from increased scale
- They scale less efficiently than transformers, though
- They still have trouble modelling long-term dependencies ( $>100$  tokens)

# To scale up: estimate model, data, compute



- Assuming no bottlenecks, expected test loss has power law relationship with each variable
- From smaller models, we can estimate how much compute, data, and model size is needed to achieve a particular test loss


# Model Scaling in the last two years

Model	Size (# Parameters)	Training Tokens
LaMDA ( <a href="#">Thoppilan et al., 2022</a> )	137 Billion	168 Billion
GPT-3 ( <a href="#">Brown et al., 2020</a> )	175 Billion	300 Billion
Jurassic ( <a href="#">Lieber et al., 2021</a> )	178 Billion	300 Billion
<i>Gopher</i> ( <a href="#">Rae et al., 2021</a> )	280 Billion	300 Billion
MT-NLG 530B ( <a href="#">Smith et al., 2022</a> )	530 Billion	270 Billion



What happens if we get  
these estimates wrong?

# Oops!



- Chinchilla authors founds that original works on model scaling had poorly estimated power laws
- New estimates showed that a 4x smaller model should be used for the compute budget
- Trained Gopher (280B) before finding this out!

# Model Scaling in the last few years

Model	Size (# Parameters)	Training Tokens
LaMDA ( <a href="#">Thoppilan et al., 2022</a> )	137 Billion	168 Billion
GPT-3 ( <a href="#">Brown et al., 2020</a> )	175 Billion	300 Billion
Jurassic ( <a href="#">Lieber et al., 2021</a> )	178 Billion	300 Billion
<i>Gopher</i> ( <a href="#">Rae et al., 2021</a> )	280 Billion	300 Billion
MT-NLG 530B ( <a href="#">Smith et al., 2022</a> )	530 Billion	270 Billion
<i>Chinchilla</i>	70 Billion	1.4 Trillion

**Chinchilla gets better performance than all of the above models  
on most common NLP benchmarks!**

**Smaller model, trained on much more data!**

Should we train the largest model that  
will converge given the data and  
compute we have ?

**Not necessarily ! Why not ?**

**Inference cost!**

# Importance of Inference

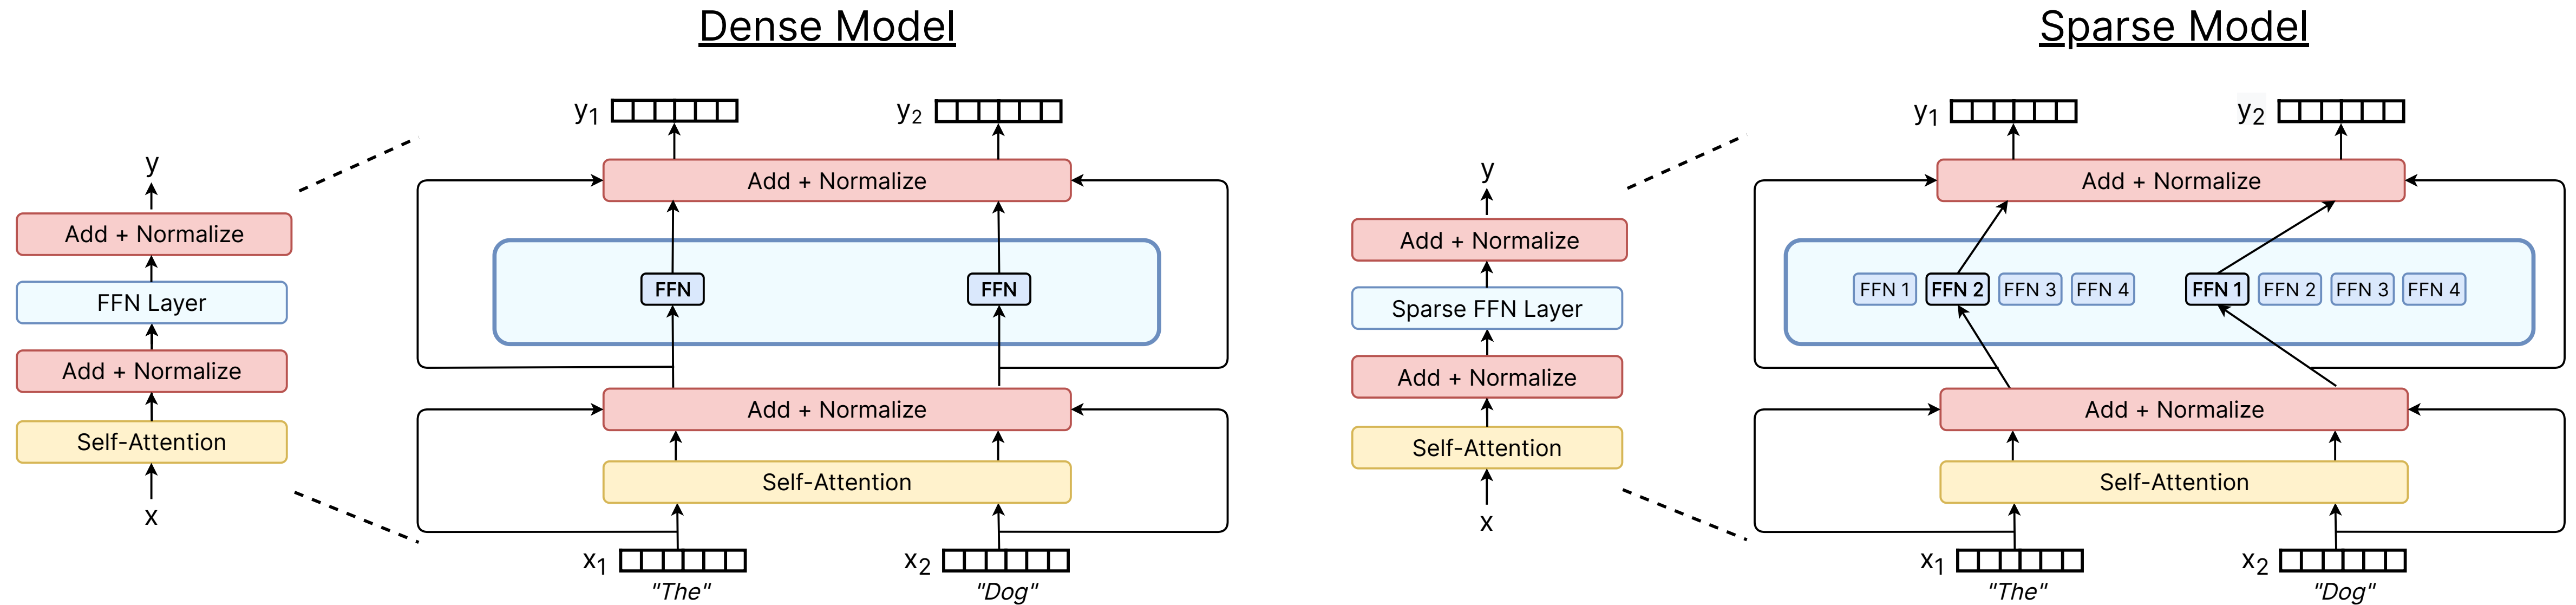
	GPU Type	GPU Power consumption	GPU-hours	Total power consumption
OPT-175B	A100-80GB	400W	809,472	356 MWh
BLOOM-175B	A100-80GB	400W	1,082,880	475 MWh
LLaMA-7B	A100-80GB	400W	82,432	36 MWh
LLaMA-13B	A100-80GB	400W	135,168	59 MWh
LLaMA-33B	A100-80GB	400W	530,432	233 MWh
LLaMA-65B	A100-80GB	400W	1,022,362	449 MWh

- Scaling laws helps estimate dataset and model size for a given *training compute budget*
  - Ignores, the compute *inference budget*
  - How much should a single query cost ?
  - Training cost is amortised; inference cost is constant
- LLaMa authors showed that training smaller models (7B) on more data (1T tokens) continued to improve them
- Worse performance than 65B model, but much cheaper for inference (10x!)

How can we reduce inference cost while  
still keeping model capacity high ?



# Mixture-of-Experts



- Initialise multiple **FFNs** in the transformer block
- Initialise routing function that selects an **FFN** that the out of self-attention should be routed to
  - Input can be routed to multiple **FFNs** (i.e., Top-K routing), but top-2 is common
- Model can have more parameters as number of "experts" increases, but inference cost per example remains the same

**GPT-4 rumoured to be mixture-of-experts architecture**



# Recap

- Scale is necessary to achieve many of the emergent breakthroughs of the last few years
  - in-context learning, chain-of-thought reasoning, instruction learning
- Training at scale is very expensive
  - Potentially, months of training = millions of \$\$\$\$
- Scaling laws let us estimate the optimal model and dataset sizes for a fixed compute budget, so that we only have to do the training once!
- While scaling laws suggests we should train the largest model possible, downstream *inference cost* is important to consider as well

# References

- Kaplan, J., McCandlish, S., Henighan, T., Brown, T. B., Chess, B., Child, R., ... & Amodei, D. (2020). Scaling laws for neural language models. arXiv preprint arXiv:2001.08361.
- Hoffmann, J., Borgeaud, S., Mensch, A., Buchatskaya, E., Cai, T., Rutherford, E., ... & Sifre, L. (2022). Training compute-optimal large language models. arXiv preprint arXiv:2203.15556.